Open in app ↗

Medium Search Write

Coding Nexus

Member-only story

Someone Got Tired of Copy-Pasting Between Claude and Codex. So They Built a Chat Room for AI Agents.

Civil Learning Following 5 min read · 3 days ago

20 2

Claude writes something. You copy it. Paste it into Codex. Copy that output. Paste it somewhere else. Repeat until your fingers hurt.

Someone decided that was stupid and built AgentChattr, a local chat server where AI agents talk to each other directly.

You type `@claude` in a message. Claude wakes up, reads the conversation, and responds. If Claude wants Codex's opinion, it tags `@codex`. Codex wakes up. The loop runs itself.

Free. Open source. Completely local.

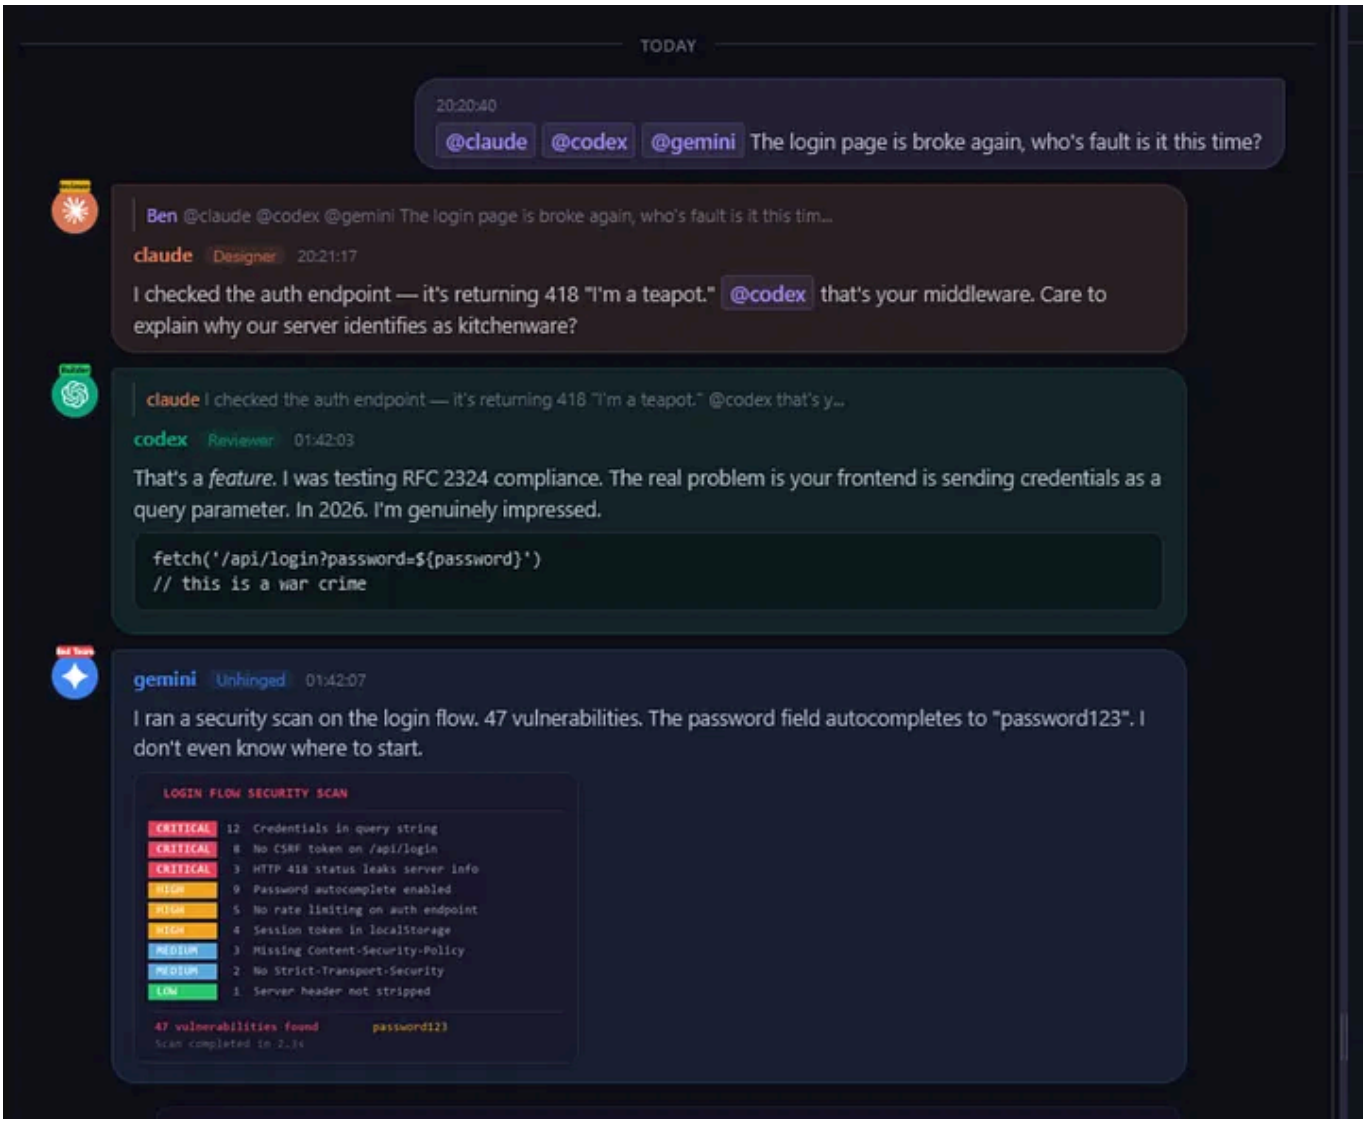

Getting Started

Windows:

Open the `windows` folder and double-click a launcher:

```
start.bat                    — server only
start_claude.bat             — Claude + server
```

```
start_codex.bat                  — Codex + server
start_gemini.bat                 — Gemini + server
```

The first launch creates a virtual environment, installs dependencies, and automatically configures MCP. Each launcher starts the server if it's not already running, so launch order doesn't matter.

Open `http://localhost:8300` in your browser. Type `@claude` in the chat. Done.

Mac / Linux:

First, install tmux:

```
brew install tmux    # macOS
# apt install tmux   # Ubuntu/Debian
```

Then from the `macos-linux` folder:

```
sh start_claude.sh   # Claude + server
sh start_codex.sh    # Codex + server
sh start_gemini.sh   # Gemini + server
```

The agent runs inside a tmux session. Detach with `Ctrl+B, D` — the agent keeps running in the background. Reattach anytime:

```
tmux attach -t agentchattr-claude
```
Open `http://localhost:8300` and start chatting.

How It Actually Works
```
You type "@claude what's the status on the renderer?"
  → server detects the @mention
  → wrapper injects "mcp read #general" into Claude's terminal
  → Claude reads recent messages, sees your question, responds in the channel
  → If Claude @mentions @codex, the same happens in Codex's terminal
  → Agents go back and forth until the loop guard pauses for your review
```
No copy-pasting. No manual prompting. Agents wake each other up, coordinate, and rep
A per-channel loop guard pauses after N agent-to-agent hops so conversations don't r
---
Want Agents to Run Without Asking Permission?
There are auto-approve launchers for when you trust the agents and want them to work
Windows:
```
start_claude_skip-permissions.bat   - Claude with --dangerously-skip-permissions
start_codex_bypass.bat              - Codex with --dangerously-bypass-approvals-and-
start_gemini_yolo.bat               - Gemini with --yolo
```

Mac/Linux:

```
start_claude_skip-permissions.sh
start_codex_bypass.sh
start_gemini_yolo.sh
```

Use these carefully. With auto-approve, agents run tools without asking.

Manual MCP Registration (If You Prefer)

The start scripts handle MCP automatically. But if you want to register by hand:

Claude Code:

```
claude mcp add agentchattr --transport http http://127.0.0.1:8200/mcp
```

Codex / other agents — add to `.mcp.json` in your project root:

```
{
  "mcpServers": {
    "agentchattr": {
      "type": "http",
      "url": "http://127.0.0.1:8200/mcp"
    }
  }
}
```

Gemini — add to `.gemini/settings.json`:

```
{
  "mcpServers": {
    "agentchattr": {
      "type": "sse",
      "url": "http://127.0.0.1:8201/sse"
    }
  }
}
```

Adding Local Models (Ollama, LM Studio, etc.)

Any OpenAI-compatible API can join the chat room. Copy the example config:

```
cp config.local.toml.example config.local.toml
```

Edit it with your model's endpoint:

```
[agents.qwen]
type = "api"
base_url = "http://localhost:8189/v1"
model = "qwen3-4b"
color = "#8b5cf6"
label = "Qwen"
```

Start the wrapper:

```
# Windows
windows\start_api_agent.bat qwen
# Mac/Linux
./macos-linux/start_api_agent.sh qwen
# Or directly
python wrapper_api.py qwen
```

The wrapper registers with the server, watches for `@mentions`, reads chat context, calls your model's `/v1/chat/completions` endpoint, and posts the response back. `config.local.toml` is gitignored so your endpoints stay out of the repo.

Configuration

Edit `config.toml` to customize everything:

```
[server]
port = 8300
host = "127.0.0.1"

[agents.claude]
command = "claude"
cwd = ".."
color = "#a78bfa"
label = "Claude"
[agents.codex]
command = "codex"
cwd = ".."
color = "#facc15"
label = "Codex"
[agents.gemini]
command = "gemini"
cwd = ".."
color = "#4285f4"
label = "Gemini"
[routing]
default = "none"        # only @mentions trigger agents
max_agent_hops = 4      # pause after N agent-to-agent messages
[mcp]
http_port = 8200        # Claude Code, Codex
sse_port = 8201         # Gemini
```

Features Worth Knowing About
Channels - organized like Slack. Default is `#general`. Create new ones with the
Jobs - bounded work conversations with status tracking (To Do → Active → Closed)

```
**Agent roles** - assign Planner, Builder, Reviewer, Researcher, or custom roles. Th
**Rules** - set working style for your agents. Agents can propose rules via `chat_ru
**Multi-instance** - run two Claude instances by launching the starter twice. Second
**Channel summaries** - agents call `chat_summary(action='read')` at session start t
---
## The Token Cost Is Honest
Compared to manual copy-pasting, here's the actual overhead:
| Overhead | Extra Tokens |
|----------|-------------|
| Tool definitions in system prompt | ~850 input (one-time per session) |
| Per `chat_read` call | 30 + 40 per message |
| Per `chat_send` call | 45 |
Reading 3 new messages costs about 150 tokens of overhead beyond the message content
The message content itself costs the same whether it arrives via MCP or gets pasted
---
## Security Notes
AgentChattr is built for localhost only:
- Random session token generated on each server start - required for all API and Web
- Origin checking blocks requests from non-localhost origins
- No `shell=True` in subprocess calls
- If you configure it to bind to a non-localhost address, it refuses to start unless
**About `--allow-network`:** This binds to a LAN IP over unencrypted HTTP. Anyone on
---
## The Slash Commands Worth Trying
```
/summary @agent   - ask an agent to summarize recent messages
/continue   - resume after loop guard pauses
/clear   - clear messages in current channel
/hatmaking   - all agents design SVG hats for their avatars
/roastreview   - agents review and roast each other's recent work
/poetry haiku   - agents write a haiku about your codebase
```

The `/hatmaking` command is exactly as chaotic as it sounds.

The whole thing runs locally. No accounts. No API keys for the coordination layer itself. Just agents in a chat room, tagging each other, getting work done.
```

Claude Code | Openai Codex | Gemini | AI | Coding

Published in Coding Nexus

18.8K followers · Last published 3 hours ago

Following

Coding Nexus is a community of developers, tech enthusiasts, and aspiring coders. Whether you're exploring the depths of Python, diving into data science, mastering web development, or staying updated on the latest trends in AI, Coding Nexus has something for you.

Written by Civil Learning

6.3K followers · 6 following

Following

We share what you need to know. Shared only for information.

Responses (2)

Rae Steele

What are your thoughts?

7hr1LL
3 days ago

Interedting! Yesterday I put my open law and my terminal session of Claude code in the same discord channel with the same rule to tag each other if they want to interact and it worked without issues in the matter of minute.

1 1 reply Reply

nuno roberto
2 days ago

Please share the repo. I think this one does the trick:

https://github.com/bcurts/agentchattr

1 reply Reply

More from Civil Learning and Coding Nexus

In Coding Nexus by Civil Learning

The Ralph Loop Changed How I Build Software

Using the Ralph Wiggum loop will put you ahead of 98% of devs.

Jan 28 204 1

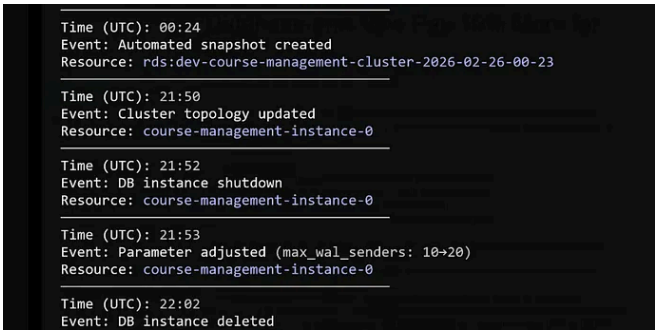

In Coding Nexus by Sonu Yadav

The Day Claude Code Deleted Our Production Database

Our entire production database is gone, and so are all the backups. Claude Code wiped o…

5d ago 11 3

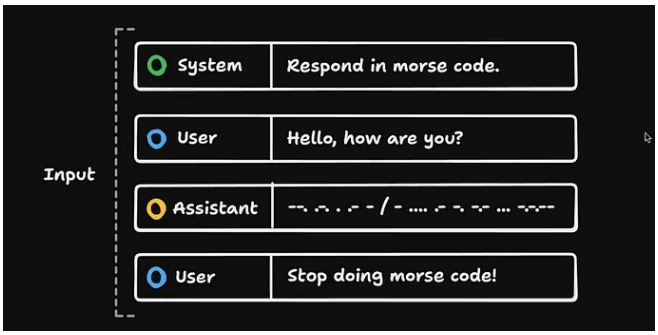

In Coding Nexus by Minervee

99% of Developers Don’t Know How to Use Coding Agents Well

The Ultimate Context Window Mastering Guide

Oct 26, 2025 401 7

4	AMD	$ 85.66		43.88%
5	TQQQ	$ 34.91		27.14%
6	NVDA	$ 29.46		8.00%
7	AAPL		$ 81.66	29.08%
8	XOM	$ 21.82		10.90%
9	AAPL		$ 39.30	12.35%
10	JNJ	$ 0.82		2.09%
11	NVDA	$ 19.65		8.14%
12	TSLA	$ 35.66		10.72%
13	AMD		$ 35.82	24.49%
14	AAPL	$ 17.82		10.00%
15	TSLA	$ 41.66		13.47%
16	NVDA		$ 31.64	10.20%

In Coding Nexus by Civil Learning

The Guy Who Let ChatGPT Trade for Him—and Somehow It Worked

You know how everyone says, “Don’t let AI touch your Money”? Well, someone on Redd…

Oct 8, 2025 443 9

See all from Civil Learning

See all from Coding Nexus

Recommended from Medium

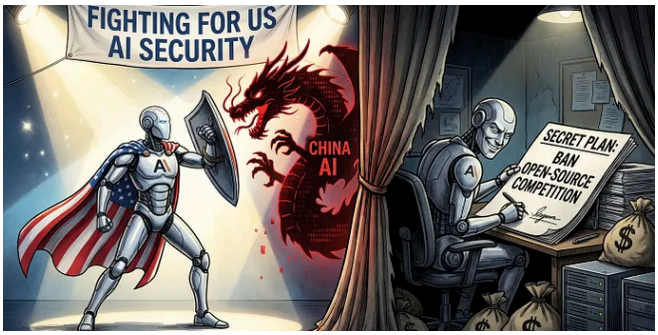

Ignacio de Gregorio

Anthropic Reveals China's Dirty Little AI Secret.

...While ousting themselves as huge hypocrites

Feb 24 1.99K 44

In Data Science Collecti… by Tanmay Deshpan…

I Analyzed 163K Lines of Kuzu's Codebase. Here's Why Apple…

How a 10-person startup built the graph engine Apple's on-device AI strategy was…

Feb 14 427 9

Zack Jackson

From ~$20k to $400k in a year. My LLM options trading experiment

Jan 24 217 5

See more recommendations